

## Case Study: Advanced Analysis of Data Analyst Job Listings in India

### Data Importing and Description

```
import pandas as pd
import numpy as np
import re
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import OneHotEncoder
```

```
# Load the dataset
df = pd.read_csv("/content/dataAnalystJobsIndia_7th_July_2024.csv")
```

```
df
```

Unnamed: 0	job_title	company	experience	min_exp	max_exp	salary	base_salary	max_salary	location	jobListed(ago)	pos
------------	-----------	---------	------------	---------	---------	--------	-------------	------------	----------	----------------	-----

JP Morgan Chase -

```
df.describe()
```

```
print(df.columns.tolist())
```

		Senior Data Analyst - IDFC FIRST Bank	5-10 Yrs	5.0	10.0	NaN	NaN	NaN	Mumbai	5.0	1
4	4	Analyst -	Bank	5-10 Yrs	5.0	10.0	NaN	NaN	NaN	Mumbai	5.0

```
print(df.dtypes)
```

```
Data  
Bangalore /  
  
missing = df.isnull().sum()  
missing_percent = (df.isnull().sum() / len(df)) * 100  
missing_summary = pd.DataFrame({  
    "Missing Count": missing,  
    "Missing %": missing_percent.round(2)  
})  
print(missing_summary)
```

	Missing	Count	Missing %
Unnamed: 0	0	0.00	
job_title	0	0.00	
Next step:	Generate code with df	0	0.00
experience	34	2.18	
min exp	34	2.18	
max exp	34	2.18	
salary	1218	78.03	
base salary	1218	78.03	
max salary	1218	78.03	
location	33	2.11	
jobListed(days ago)	0	0.00	
postedIn	0	0.00	
rating	437	27.99	
reviews count	437	27.99	
details	33	2.11	
salary data provide by	1218	78.03	

```
print(df.describe(include="number"))
```

	Unnamed: 0	min exp	max exp	base salary	max salary	\
count	1561.0000	1527.00000	1527.00000	3.430000e+02	3.430000e+02	
mean	780.0000	3.131631	6.555337	5.649980e+05	1.859911e+06	
std	450.7662	2.191749	3.199266	6.992511e+05	2.018750e+06	
min	0.0000	0.000000	1.000000	9.400000e+03	7.000000e+04	
25%	390.0000	2.000000	1.000000	1.550000e+05	7.000000e+05	

```
50%    780.0000    3.000000    6.000000  3.600000e+05  1.330000e+06
75%   1170.0000    4.000000    8.000000  7.000000e+05  2.458500e+06
max   1560.0000   15.000000   30.000000  9.000000e+06  2.250000e+07

      jobListed(days ago)      rating  reviews count
count          1561.000000  1124.000000  1124.000000
mean           223.814222    3.799377   959.846975
std            650.077316    0.740237  4198.153851
min            0.000000    1.000000   1.000000
25%           16.000000   3.500000   7.000000
50%           30.000000   3.900000   30.000000
75%           42.000000   4.200000  242.500000
max          19710.000000   5.000000  41400.000000
```

```
print(df.describe(include="object"))
```

	job_title	company	experience	salary	\
count	1561	1561	1527	343	
unique	867	1073	95	255	
top	Data Analyst	Diverse	Lynx	2-5 Yrs	₹ 10 - 20L/yr
freq	426		24	109	8

	location	postedIn	details	salary	data provide by
count	1528	1561	1528	343	
unique	194	5	1257	2	
top	Bangalore / Bengaluru	Naukri	Data Analyst	Salary Listed by Company	
freq	336	1236	20	177	

```
print("Number of duplicate rows:", df.duplicated().sum())
```

```
Number of duplicate rows: 0
```

## ▼ 1. DATA CLEANING AND PREPROCESSING

```
# Clean salary columns
def extract_salary_number(x):
    if pd.isna(x):
        return np.nan
    nums = re.findall(r"\d+\.\?\d*", str(x).replace(",", ""))
    return float(nums[0]) if nums else np.nan
```

```
df["base_salary_num"] = df["base salary"].apply(extract_salary_number)
df["max_salary_num"] = df["max salary"].apply(extract_salary_number)
```

```
# Fixing column name
if "jobListed(days ago" in df.columns:
    df.rename(columns={"jobListed(days ago": "jobListed(days ago)"}, inplace=True)
```

```
# Cleaning experience values
def clean_experience(x):
    if pd.isna(x):
        return np.nan
    x = str(x)
    match = re.findall(r"(\d+)", x)
    if len(match) == 1:
        return int(match[0])
    elif len(match) == 2:
        return (int(match[0]) + int(match[1])) / 2
    return np.nan
```

```
df["exp_span"] = df["experience"].apply(clean_experience)
```

```
# Cleaning experience column (e.g., "4-8 Yrs" → 4, 8)
def clean_exp(x):
    nums = re.findall(r"\d+", str(x))
    if len(nums) == 2:
        return int(nums[0]), int(nums[1])
    return np.nan, np.nan

df["min_exp_clean"], df["max_exp_clean"] = zip(*df["experience"].map(clean_exp))
df["exp_mid"] = df[["min_exp_clean", "max_exp_clean"]].mean(axis=1)
```


```
# Cleaning reviews count
df["reviews_count_num"] = df["reviews count"].apply(lambda x: extract_salary_number(x))
```

```
# Standardize location  
df["location_clean"] = df["location"].str.split(",").str[0].str.strip()
```


```
# Removeing rows where essential values are missing  
df = df.dropna(subset=["rating", "reviews_count_num", "exp_mid"])
```

## ▼ 2. EXPLORATORY DATA ANALYSIS


```
# Rating distribution  
plt.figure(figsize=(8,5))  
sns.histplot(df["rating"], kde=True)  
plt.title("Distribution of Job Ratings")  
plt.show()
```




```
# Reviews count distribution  
plt.figure(figsize=(8,5))  
sns.histplot(df["reviews_count_num"], kde=True)  
plt.title("Distribution of Review Counts")  
plt.show()
```



```
# Jobs by location
plt.figure(figsize=(17,7))
df["location_clean"].value_counts().plot(kind="bar")
plt.title("Job Listings by Location")
plt.xlabel("Location")
plt.ylabel("Count")
plt.show()
```



```
# Experience mid value distribution
plt.figure(figsize=(8,5))
sns.histplot(df["exp_mid"], kde=True)
plt.title("Distribution of Required Experience")
plt.show()
```



### ▼ 3. STATISTICAL ANALYSIS

```
# Correlation tests
clean_corr = df[["base salary", "reviews count"]].dropna()

if len(clean_corr) >= 2:
    pearson_r, pearson_p = stats.pearsonr(
        clean_corr["base salary"], clean_corr["reviews count"])
)
    spearman_rho, spearman_p = stats.spearmanr(
        clean_corr["base salary"], clean_corr["reviews count"])
)

print("\nPearson:", pearson_r, pearson_p)
print("Spearman:", spearman_rho, spearman_p)
```

Pearson: -0.056602251561984986 0.5240390644700293  
Spearman: 0.09986305090907503 0.2601639763732011

```
# Comparing ratings on Naukri vs iimjobs
naukri = df[df["postedIn"] == "Naukri"]["rating"]
iim = df[df["postedIn"] == "iimjobs"]["rating"]
```

```
print("\nT-test: Naukri vs iimjobs ratings")
print(stats.ttest_ind(naukri, iim, nan_policy="omit"))
```

T-test: Naukri vs iimjobs ratings  
TtestResult(statistic=np.float64(0.00047026436500566894), pvalue=np.float64(0.9996248840803188), df=np.float64(931.0))

```
# Correlation (ratings vs reviews count)
print("\nCorrelation (ratings vs reviews count)")
print(stats.pearsonr(df["rating"], df["reviews count num"]))
```

Correlation (ratings vs reviews count)  
PearsonRResult(statistic=np.float64(0.019944431027834907), pvalue=np.float64(0.5095167744732034))

## 4. MACHINE LEARNING MODEL

```
# Keeping only rows where salary exists  
ml df = df.dropna(subset=["base salary num"])
```

```
if len(ml_df) < 5:  
    print("\nX Not enough salary data to train ML model (need ≥ 5 rows).")
```

```

else:
    # Selecting features
    X = ml_df[["exp_mid", "rating", "reviews_count_num", "jobListed(days ago)"]]
    y = ml_df["base_salary_num"]

    # Handle remaining NaN
    X = X.fillna(X.mean())

    # Train-test splitting
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.3, random_state=42
    )

    model = RandomForestRegressor(random_state=42)
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

```

```

print("\nMODEL RESULTS")
print("MAE:", mean_absolute_error(y_test, y_pred))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
print("R^2:", r2_score(y_test, y_pred))

```

MODEL RESULTS  
MAE: 293894.0  
RMSE: 523202.39631273155  
R<sup>2</sup>: 0.06289312512706535

```

# Feature Importance
feat_imp = pd.DataFrame({
    "Feature": X.columns,
    "Importance": model.feature_importances_
}).sort_values(by="Importance", ascending=False)

print("\nFEATURE IMPORTANCE")
print(feat_imp)

```


FEATURE IMPORTANCE

	Feature	Importance
0	exp_mid	0.469268
2	reviews_count_num	0.210463
3	jobListed(days ago)	0.183646
1	rating	0.136624

```

# Plot feature importance
plt.figure(figsize=(7,5))
sns.barplot(data=feat_imp, x="Importance", y="Feature")
plt.title("Feature Importance in Salary Prediction")
plt.show()

```



```
# ANALYSIS BY MD RIYAZ UDDIN.
```

**ANALYSIS BY MD RIYAZ UDDIN.**

**THANK YOU...**